

Module 5:

Fine-Tuning LLMs: Unleashing Their Power

Content

- ❖ Introduction to LLM fine-tuning
- ❖ Stages in the Development of Modern LLMs
- ❖ Supervised Fine-tuning, Preference Fine-tuning (RLHF, DPO)
- ❖ Fine-Tuning With Limited Computing Resources (QLORA, Quantization)
- ❖ Hands-on experience

Introduction to fine-tuning LLMs

Fine-tuning: process of turning general-purpose models into specialized models.

joint pain, skin rash, and
sun sensitivity



Base LLM



These symptoms may be
related to inflammation

Fine-tuning: process of turning general-purpose models into specialized models.

joint pain, skin rash, and
sun sensitivity



Base LLM



These symptoms may be
related to inflammation



Fine-tuning



Fine-tuning: process of turning general-purpose models into specialized models.

joint pain, skin rash, and
sun sensitivity



Base LLM



These symptoms may be
related to inflammation

joint pain, skin rash, and
sun sensitivity



Base LLM



These symptoms suggests potential
autoimmune involvement. Conditions
like lupus often cause these symptoms
due to photosensitivity.

Fine-tuning

Allergy
data

Fine-tuning LLMs

Fine-tuning involves **adjusting the parameters** of a pre-trained LLM using a smaller, task-specific dataset.

The goal is to **customize the model** for specific language patterns and vocabulary related to a particular task.

Benefits of fine-tuning

- **Specificity and Relevance:** Fine-tuning ensures LLMs understand industry-specific terms and generate relevant content.
- **Improved Accuracy:** Domain-specific fine-tuning enhances precision, aligning model outputs with expectations.
- **Customized Interactions:** Tailoring LLM responses maintains brand consistency and user experience.

Benefits of fine-tuning

- **Data Privacy and Security:** Fine-tuning controls exposure to sensitive data, preventing inadvertent leaks.
- **Addressing Rare Scenarios:** Fine-tuning optimally handles unique business challenges.

Stages in the Development of Modern Foundation Models

Development of Modern Foundation Models

Base LLM

The benefits of regular exercise

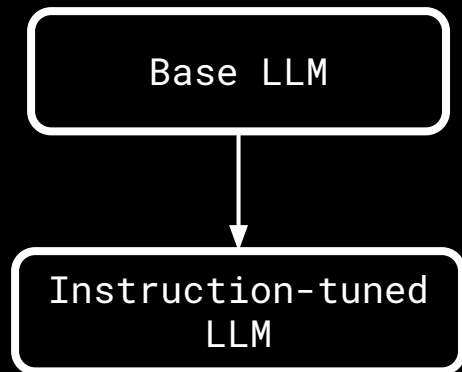


Pretrained LLM

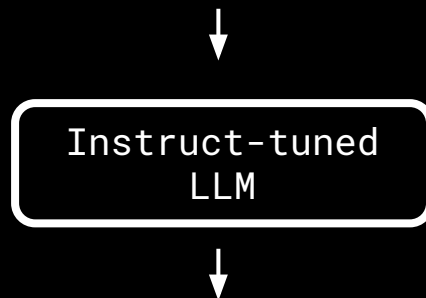


The benefits of regular exercise are well-documented. Exercise can improve mood, boost energy,

Development of Modern Foundation Models

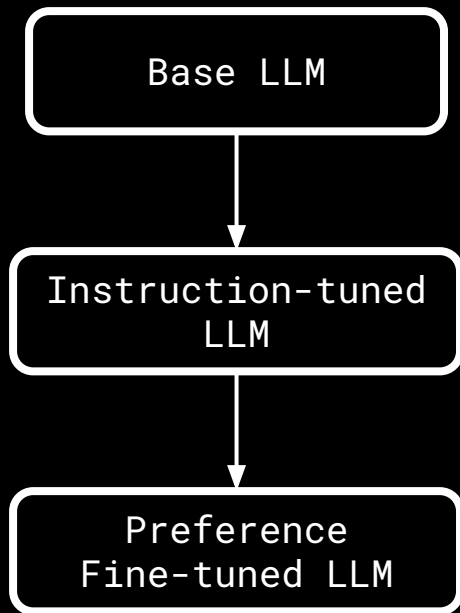


The benefits of regular exercise

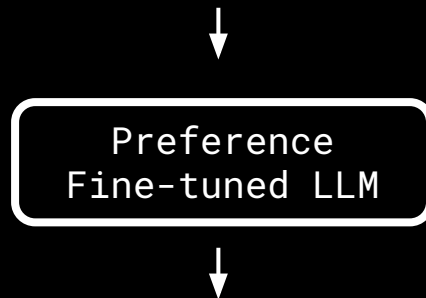


Regular exercise offers numerous benefits for both physical and mental health. Here are some of the main advantages: Weight Management: Regular exercise helps in maintaining a healthy weight by burning calories and boosting metabolism improved Mood: ns

Development of Modern Foundation Models



The benefits of regular exercise



Regular exercise offers a multitude of advantages: Weight management: Burn calories and maintain a healthy weight. Disease prevention: Reduce risks of heart disease, stroke, diabetes, and certain cancers.

Mood enhancement: Boost happiness..

Pretraining LLM

Self-Supervised Setting:

The model predicts the next word in a given context.

Loss Function:

Cross Entropy Loss

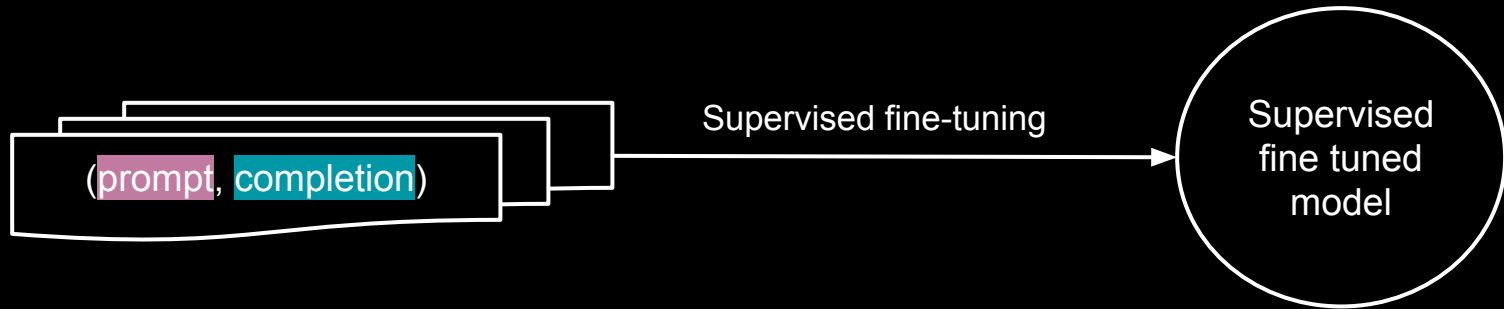
Data:

LLMs are pre-trained on vast corpora from the internet.

This diverse data source enables the model to capture general language patterns and understanding

Supervised fine-tuning

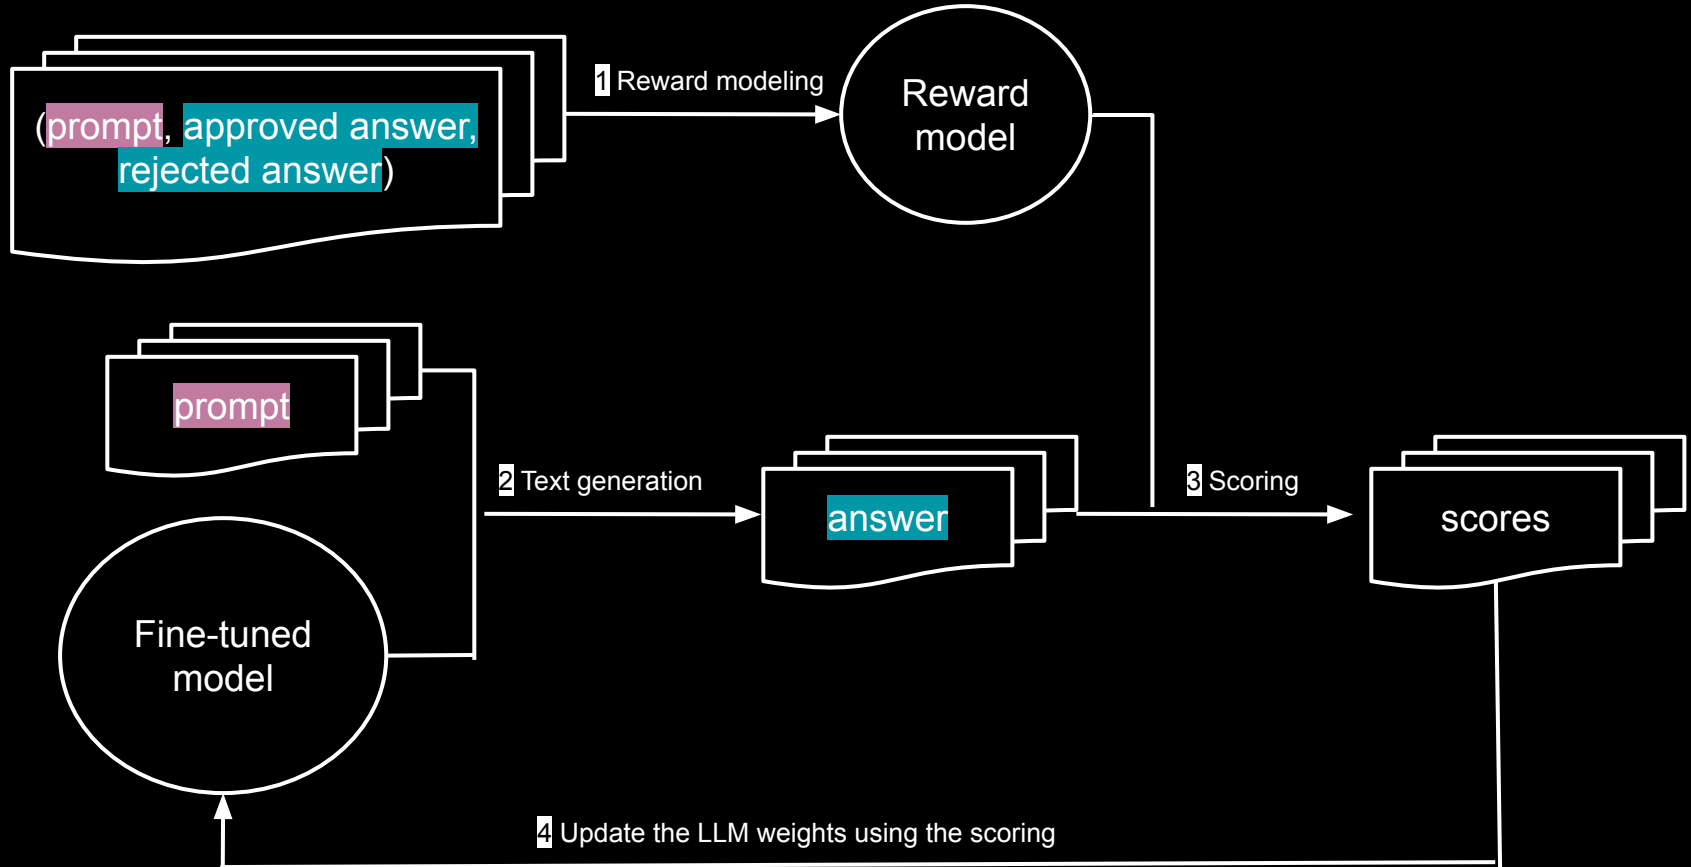
- **Purpose:** Enable an LLM to perform specific tasks
- **Requirements:** Large dataset of prompts and correct answers for the task
- **Process:** maximise the likelihood of tokens in the answers
- Cross Entropy Loss



Preference fine-tuning

- **Purpose:** adjust the LLM to better reflect the preferences from a comparison dataset
- **Relevance:**
 - useful for critiquing LLM-generated answers
 - Faster and easier than manually writing adequate answers.
- **Comparison dataset:** contains prompts, approved answers and rejected answers
- **Fine-tuning methods**
 - Reinforcement Learning with Human Feedback (RLHF)
 - Direct Preference Optimization (DPO)

Reinforcement Learning with Human Feedback



Reinforcement Learning with Human Feedback

$$\max_{\pi_{\theta}} \mathbb{E}_{x \sim \mathcal{D}, y \sim \pi_{\theta}(y|x)} [r_{\phi}(x, y)] - \beta \mathbb{D}_{\text{KL}} [\pi_{\theta}(y | x) || \pi_{\text{ref}}(y | x)]$$

maximise
rewards

use KL-divergence penalty to prevent
reward hacking (controlled by β)

Reinforcement Learning with Human Feedback

$$\max_{\pi_{\theta}} \mathbb{E}_{x \sim \mathcal{D}, y \sim \pi_{\theta}(y|x)} [r_{\phi}(x, y)] - \beta \mathbb{D}_{\text{KL}} [\pi_{\theta}(y | x) || \pi_{\text{ref}}(y | x)]$$

maximise
rewards

use KL-divergence penalty to prevent
reward hacking (controlled by β)

- Various challenges

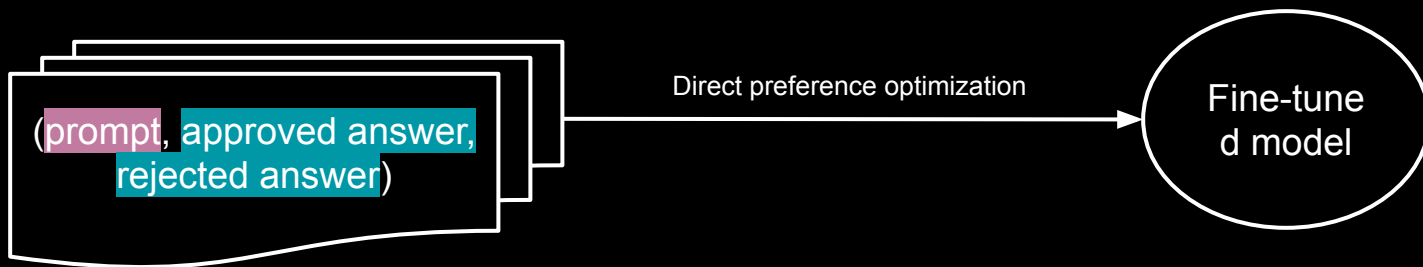
RL notoriously unstable, many hyperparameters

Need a separate RM \Rightarrow 3 LLMs to jungle

Direct Preference Optimization

It solves the same problem by minimizing a training loss directly based on the preference data (without reward modeling or reinforcement learning)

Simpler and more stable alternative



Direct Preference Optimization

$$\max_{\pi} \mathbb{E}_{(x, y_w, y_l) \sim \mathcal{D}} \log \sigma \left(\beta \log \frac{\pi(y_w|x)}{\pi_{\text{ref}}(y_w|x)} - \beta \log \frac{\pi(y_l|x)}{\pi_{\text{ref}}(y_l|x)} \right)$$

Rafailov and al. (2023)

Direct Preference Optimization

$$\max_{\pi} \mathbb{E}_{(x, y_w, y_l) \sim \mathcal{D}}$$



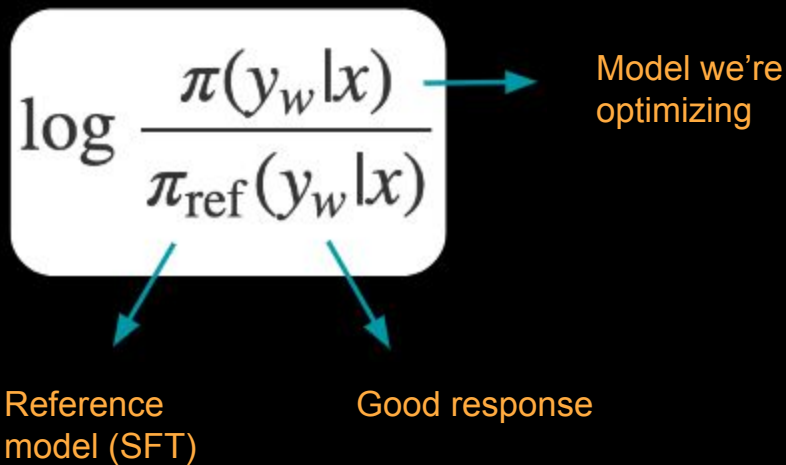
Good

response

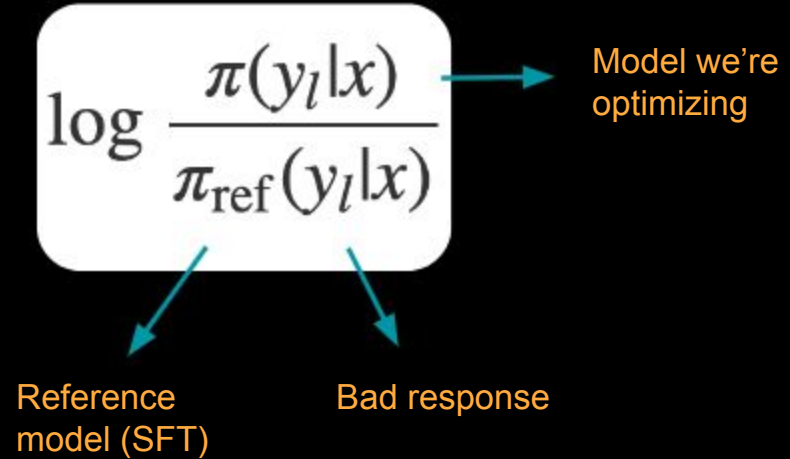
bad

response

Direct Preference Optimization



Direct Preference Optimization



Direct Preference Optimization



$$\max_{\pi} \mathbb{E}_{(x, y_w, y_l) \sim \mathcal{D}} \log \sigma \left(\beta \log \frac{\pi(y_w|x)}{\pi_{\text{ref}}(y_w|x)} - \beta \log \frac{\pi(y_l|x)}{\pi_{\text{ref}}(y_l|x)} \right)$$

```
import torch.nn.functional as F
```

```
def dpo_loss(pi_logps, ref_logps, yw_idx, yl_idx, beta):  
    pi_yw_logps, pi_yl_logps = pi_logps[yw_idx], pi_logps[yl_idx]  
    ref_yw_logps, ref_yl_logps = ref_logps[yw_idx], ref_logps[yl_idx]  
    pi_logratios = pi_yw_logps - pi_yl_logps  
    ref_logratios = ref_yw_logps - ref_yl_logps  
    losses = -F.logsigmoid(beta * (pi_logratios - ref_logratios))  
    rewards = beta * (pi_logps - ref_logps).detach()  
    return losses, rewards
```

Algorithm

- Sample good/bad response
- Run pairs through 2 models (active and reference)
- Backpropagation

Efficient fine-tuning

Vanilla fine-tuning of LLMs

- It refers to the process of adjusting all parameters of a pre-trained model when further training it on a new dataset.
- Challenges
 - Parameter Count
 - Gradient computation and storing (memory-intensive)
 - Computationally expensive and time-consuming.

Parameters-Efficient Fine-Tuning (PEFT)

PEFT Paradigm:

Address the challenges of vanilla fine-tuning.

Target Parameters:

PEFT identifies specific layers or parameters that significantly impact the task performance.

Freezing Others:

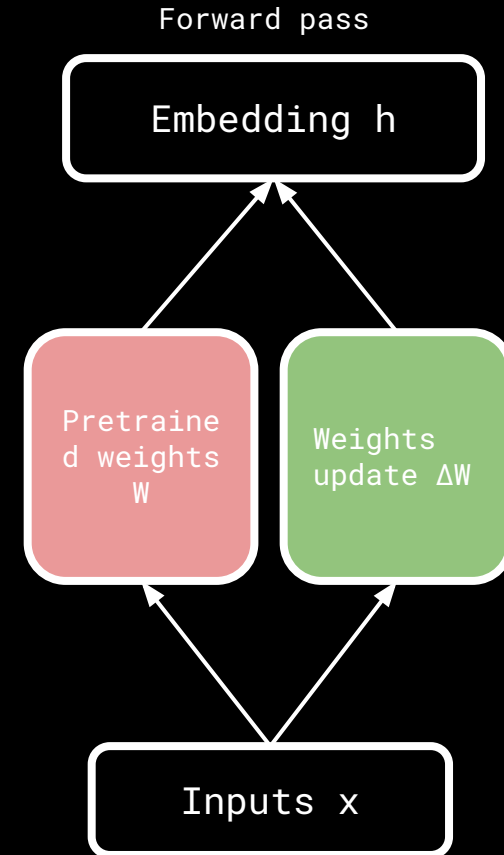
The rest of the learned parameters (non-target parameters) are frozen.

Low Rank Adaptation : LoRA

Fine-tuning: leverages general knowledge
From a pretrained model

- 405 B parameters for Llama 3.1

Fine-tuning \Rightarrow **memory intensive**

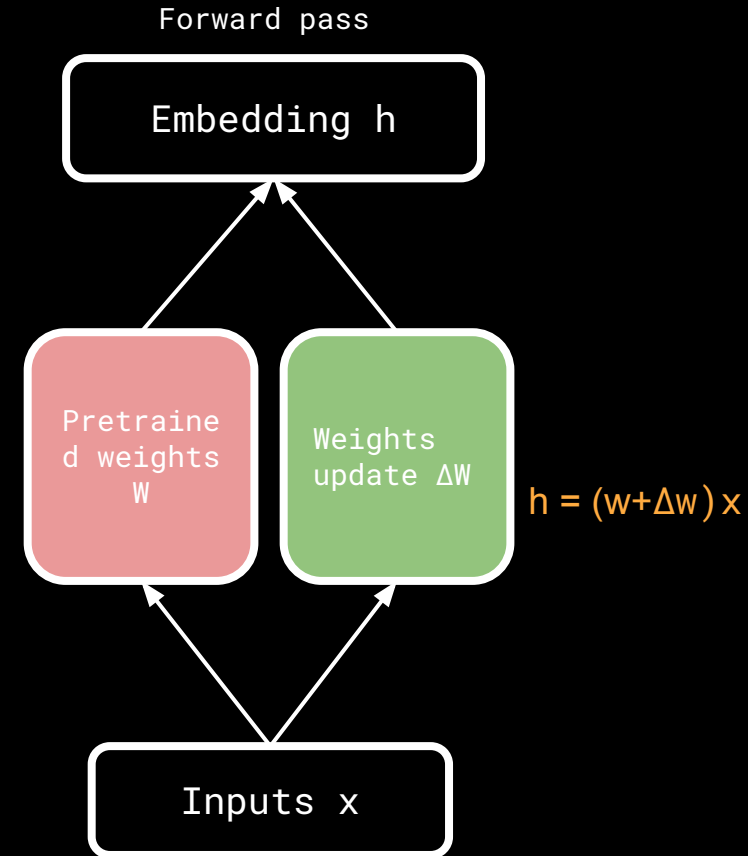


Low Rank Adaptation : LoRA

Fine-tuning: leverages general knowledge
From a pretrained model

- 405 B parameters for Llama 3.1

Fine-tuning \Rightarrow **memory intensive**



Low Rank Adaptation : LORA

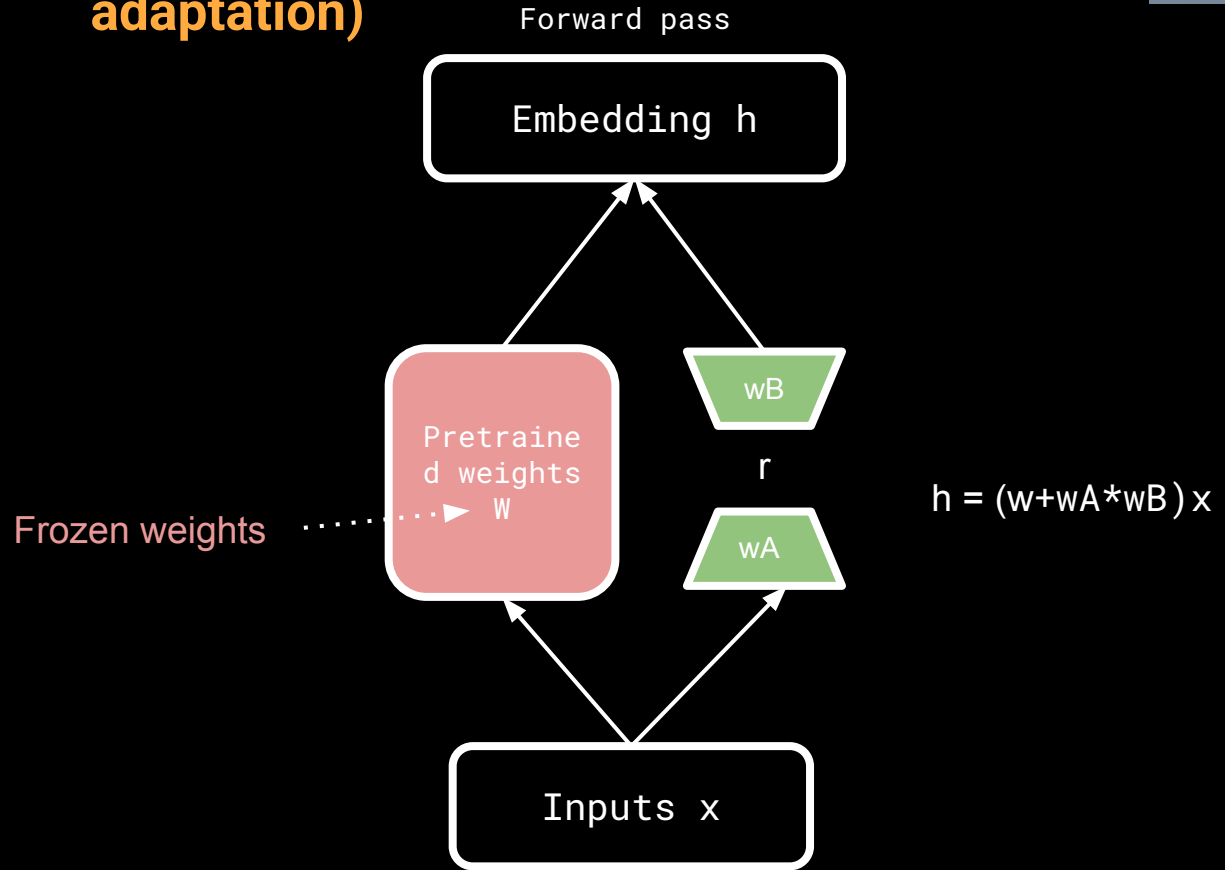
Hu et al. (2021)

- During fine-tuning, ΔW has a low rank and can be decomposed as

$$\begin{array}{c} \Delta W \\ A \left[\begin{array}{c} \\ \\ \end{array} \right] B \end{array} = \begin{array}{c} WA \\ A \left[\begin{array}{c} \\ \\ \end{array} \right] \begin{array}{c} r \\ \end{array} \end{array} \times \begin{array}{c} WB \\ \left[\begin{array}{c} \\ \\ \end{array} \right] B \end{array}$$

- r is hyper-parameter that we need to tune

Fine-Tuning With Limited Computing Resources: LORA (Low rank adaptation)



Quantization

Floating-Point Representation

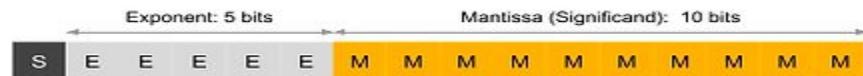
fp32: Single-precision IEEE Floating Point Format

Range: $\sim 1e^{-38}$ to $\sim 3e^{38}$



fp16: Half-precision IEEE Floating Point Format

Range: $\sim 5.96e^{-8}$ to 65504



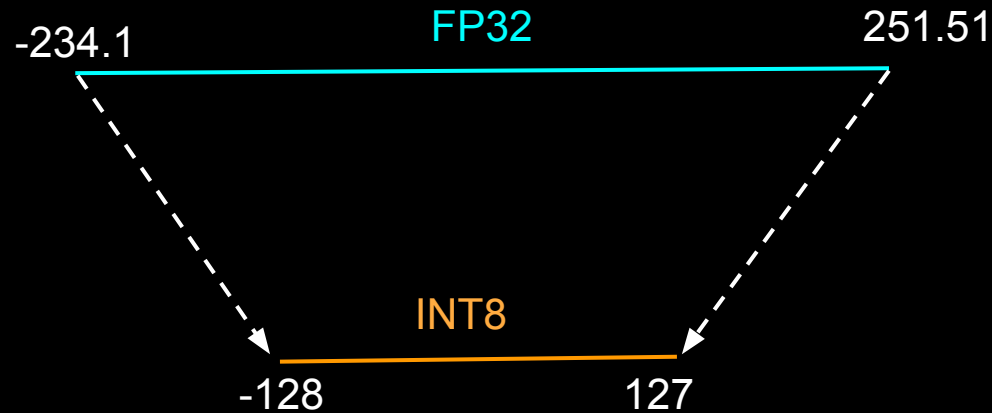
bfloat16: Brain Floating Point Format

Range: $\sim 1e^{-38}$ to $\sim 3e^{38}$



Quantization

Quantization refers to the process of mapping input values from a large set (often continuous) to output values in a smaller set, often with a finite number of elements.



Quantization

Zero point quantization

Let consider

2.8912 -0.1244 4.1234 1.9876 -1.4567

Quantization

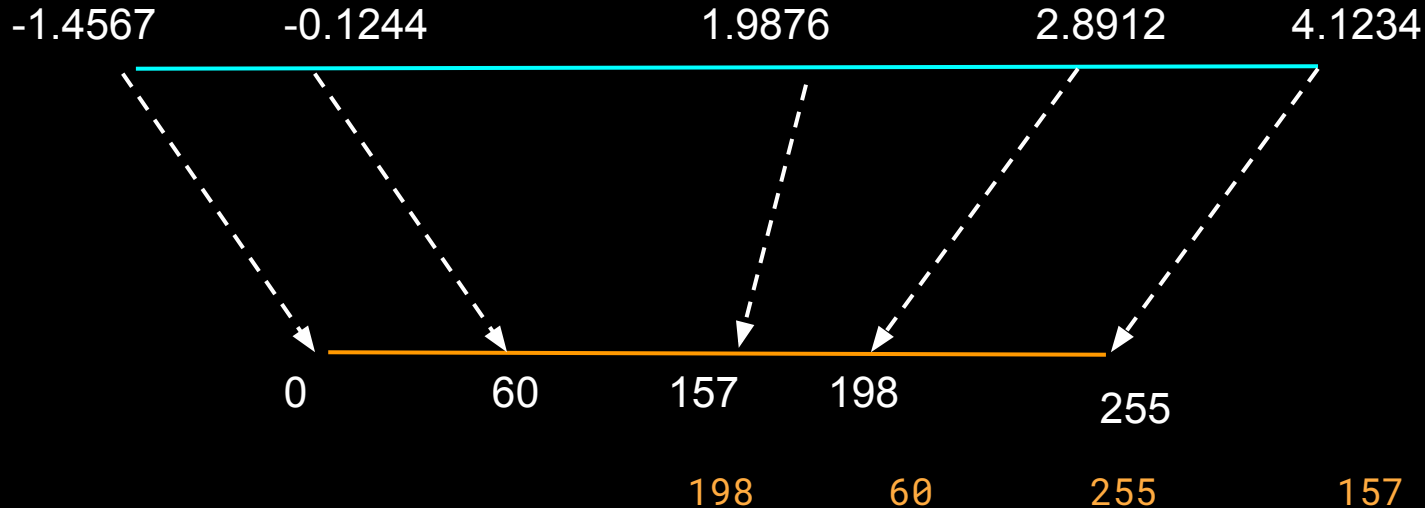
Zero point quantization

Let consider

2.8912 -0.1244 4.1234 1.9876 -1.4567

Max

min



Quantized LoRA: QLoRA



Dettmers et al (2023)

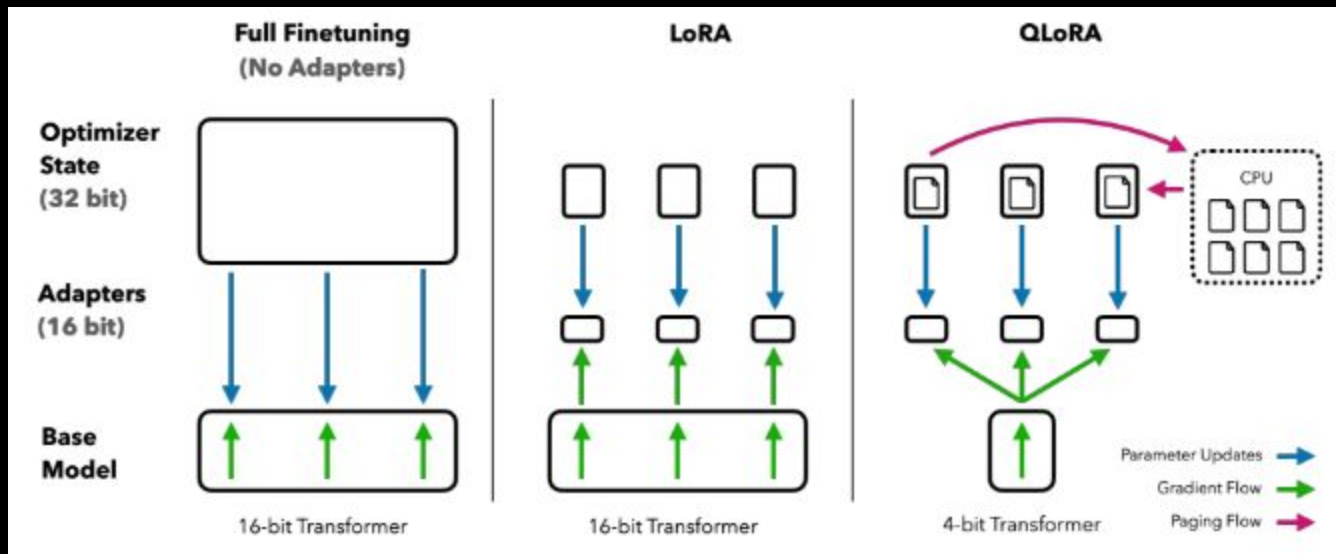
QLoRA: efficient fine-tuning approach designed to reduce memory usage during the fine-tuning of LLMs.

It enables fine-tuning of a **65B parameter model** on a **single 48GB GPU** while maintaining full 16-bit fine-tuning task performance.

The approach involves back-propagating gradients through a frozen, 4-bit quantized pretrained language model into Low Rank Adapters (LoRA).

Fine-Tuning With Limited Computing Resources: QLoRA (Quantized LORA)

Dettmers et al (2023)





AI TRAINING

Fine-Tuning Best Practices

- Clearly Define Your Task

Foundational Step: Begin by defining your specific task.

Focus and Direction: Clear task definition channels the model's capabilities toward a specific goal.

Performance Benchmarks: Set measurable benchmarks for evaluating model performance.

- Leveraging Pre-Trained Models

Efficiency and Understanding: Pre-training captures general language understanding.

Model Architecture Matters: Choose the right architecture (e.g., MoE, MoT) for effective fine-tuning.

Set Hyperparameters

Tunable Variables: Hyperparameters (e.g., learning rate, batch size, weight decay) impact model training.

Optimal Configuration: Experiment to find the best hyperparameter values for your specific task.

Iterative Refinement: Continuously evaluate and adjust hyperparameters during fine-tuning.

Evaluate Model Performance

Unbiased Assessment: Evaluate the fine-tuned model on a separate test set.

Generalization: Assess how well the model performs on unseen data.

Refinement Opportunity: If performance can be improved, consider further iterations.